

# M4: Shared Memory Programming with OpenMP


1

- Shared Memory Programming with OpenMP: A Standard for Directive Based Parallel Programming. (6-Hrs )
  - Writing programs that use OpenMP.
  - Using OpenMP to parallelize many serial for loops with only small changes to the source code.
  - Task parallelism.
  - Explicit thread synchronization.
  - Standard problems in shared-memory programming

# Writing programs that use OpenMP.

- An API for shared-memory parallel programming.
  - MP = multiprocessing
  - Designed for systems in which each thread or process can potentially have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to main memory.

• Shared Memory Systems  
Conceptual model  
Real Hardware model



- OpenMP is a multi-threading, shared address model.
  - Threads communicate by sharing variables.
  - Unintended sharing of data causes race conditions
  - To control race conditions: – Use synchronization to protect data conflicts.

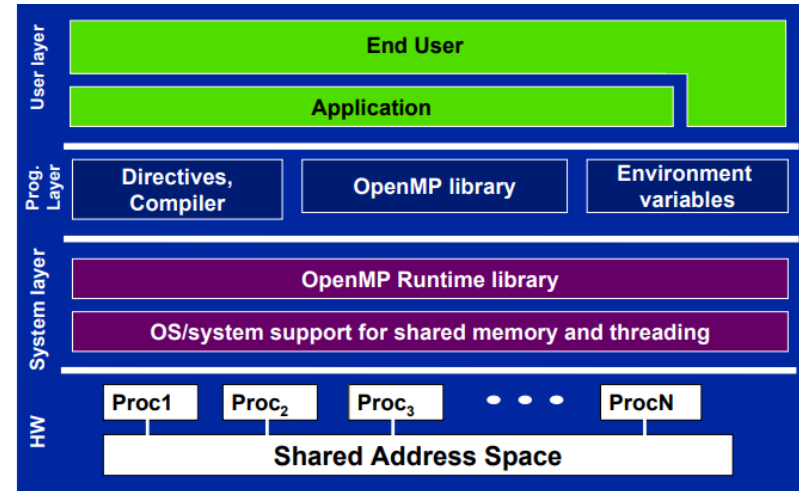
# OpenMP core syntax

4

- Most of the constructs in OpenMP are compiler directives.
  - `#pragma omp construct [clause [clause]...]`
- Function prototypes and types in the file: `#include<omp.h>`.
- constructs apply to a “structured block”

```
#pragma omp parallel num_threads(4)
```

Request a certain number of threads



# OpenMp pragmas

`# pragma omp parallel`

Most basic parallel directive

`# pragma omp parallel num_threads ( thread_count )`

It allows the programmer to specify the number of threads that should execute the following block

**Implicit barrier** : thread that has completed the block of code will wait for all the other threads in the team to complete the block

# Pragmas

6

- **Pragmas**

- Special preprocessor instructions.
- Typically added to a system to allow behaviors that aren't part of the basic C specification.
- Compilers that don't support the pragmas ignore them

```
#pragma omp parallel  
{  
    //Parallel region code  
}
```

#pragma

```
gcc -g -Wall -fopenmp -o omp_hello omp_hello . C  
/omp_hello 4
```

export OMP\_NUM\_THREADS=5

# Exercise 1:

7

```
#include <stdio.h>

int main( int ac, char **av)
{
    //a and shared between all thread (race cond)
    #pragma omp parallel
    {
        int a;
        int b;
        a=omp_get_num_threads();
        b=omp_get_thread_num();
        printf("Total No of Threads %d",a);
        printf("Hello World!!! with thread ID %d \n ",b);
    }

    return 0;
}
```

original thread and the new threads — is called a team, the original thread is called the master, and the additional threads are called slaves.

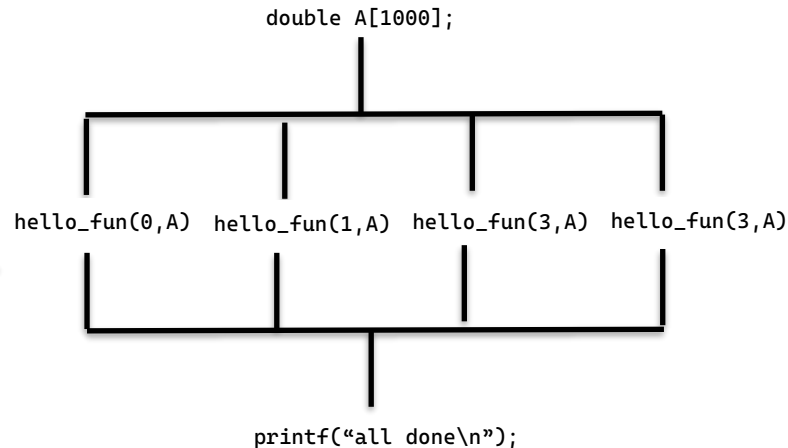
```
Total No of Threads 8
Total No of Threads 8
Total No of Threads 8
Hello World!!! with thread ID 1
Total No of Threads 8
Hello World!!! with thread ID 7
Hello World!!! with thread ID 0
Hello World!!! with thread ID 4
Total No of Threads 8
Total No of Threads 8
Total No of Threads 8
Hello World!!! with thread ID 6
Hello World!!! with thread ID 2
Total No of Threads 8
Hello World!!! with thread ID 3
Total No of Threads 8
Hello World!!! with thread ID 5
```

```
Total No of Threads 8
Total No of Threads 8
Hello World!!! with thread ID 5
Total No of Threads 8
Total No of Threads 8
Hello World!!! with thread ID 7
Hello World!!! with thread ID 4
Total No of Threads 8
Total No of Threads 8
Hello World!!! with thread ID 3
Hello World!!! with thread ID 2
Hello World!!! with thread ID 0
Total No of Threads 8
Hello World!!! with thread ID 1
Total No of Threads 8
Hello World!!! with thread ID 6
```

# Thread Creation: Parallel Regions example

8

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{ int ID = omp_get_thread_num();  
  hello_fun(ID, A);  
}  
printf("all done\n");
```





1. passing the number of threads at the command line

# Parallel construct

10

```
#pragma omp parallel [clause [, clause] ...]  
    structured-block
```

**The following clauses apply:**

- if
- num\_threads
- shared, private, firstprivate, default
- reduction
- copyin
- proc\_bind

# Parallel construct

11

- `num_threads` clause
  - Specifies how many threads should execute the region
  - Syntax: “`num_threads (scalar-logical-expression)`”
- `if` clause
  - Conditional parallel execution
  - Avoid parallelization overhead if little work to be parallelized
  - Syntax: “`if (scalar-logical-expression)`”
  - If the logical expression evaluates to true: execute in parallel

Example:

```
int main( ... )
{
    [...]
    #pragma omp parallel num_threads (nths)
    {
        [...]
    }
    [...]
}
```

```
int main( ... )
{
    [...]
    #pragma omp parallel if (n > 1000)
    {
        [...]
    }
    [...]
}
```

# Thread synchronization

12

- High level synchronization: –
  - Critical
  - atomic
  - barrier
  - Ordered
- Low level synchronization
  - – flush
  - – locks

# Explicit thread synchronization.

13

- High level synchronization: –
  - Critical
  - atomic
  - barrier
  - Ordered
- Low level synchronization
  - – flush
  - – locks

# Synchronization

14

- Critical : Mutual exclusion: Only one thread at a time can enter a critical region

```
float res;
#pragma omp parallel
{
    float B; int i, id, nthrds;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    for(i=id; i<niters; i+=nthrds)
    {
        B = big_job(i);
        #pragma omp critical
            consume (B, res);
    }
}
```

# Synchronization

15

- Atomic provides mutual exclusion but only applies to the update of a memory location (the update of X in the following example)

```
#pragma omp parallel
{
  double tmp, B;
  B = func1();
  tmp = func2(B);
  #pragma omp atomic
  X += func2(B);
}
```

**Atomic only protects the read/update of X**

- The number of threads in a parallel region is determined by the following factors, in order of precedence.
  - Evaluation of the if clause
  - Setting of the num\_threads() clause
  - Use of the omp\_set\_num\_threads() library function
  - Setting of the OMP\_NUM\_THREADS environment variable



# Create threads with the parallel construct

17

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
int main()
{
    int nthreads, tid;
    #pragma omp parallel num_threads(4) private(tid)
    {
        tid = omp_get_thread_num();
        printf("Hello world from (%d)\n", tid);
        if(tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("number of threads = %d\n", nthreads);
        }
    }
}
```

Write openMP program to fork a group of threads with each thread having a private variable to store thread number

```
#include<stdlib.h>
#include <stdio.h>
#include "omp.h"
int main()
{  int nthreads, A[100] , tid; //
    omp_set_num_threads(4);
    #pragma omp parallel private (tid)
        { tid = omp_get_thread_num();
          Fun1(tid, A);
        }
}
```

# Thread termination

19

- By default, worksharing for loops end with an implicit barrier
  - **nowait**: If specified, threads do not synchronize at the end of the parallel loop  
`#pragma omp parallel nowait`
  - **ordered**: specifies that the iteration of the loop must be executed as they would be in serial program.
  - **collapse**: specifies how many loops in a nested loop should be collapsed into one large iteration space and divided according to the schedule clause.
  - The sequential execution of the iteration in all associated loops determines the order of the iterations in the collapsed iteration space

Barrier: each thread for s waits till all threads arrive  
`#pragma omp barrier`

No implicit barrier due to nowait.

`#pragma omp for nowait`

# OpenMP programming

20

- Parallel Construct
- Work-Sharing Constructs
  - Loop Construct
  - Sections Construct
  - Single Construct
- Synchronization constructs
  - Barrier Construct
  - Critical Construct
  - Atomic Construct
- Data clauses
  - shared, private, Lastprivate, firstprivate, default

- Parallel construct

- Used to specify the computations that should be executed in parallel.

**#pragma omp parallel** [*clause*[[,] *clause*]. . . ]

..... code block .....

- The work of the region is replicated for every thread.

# Work-Sharing Construct

- Work-sharing is to split up pathways through the code between threads within a team
  - Loop construct
    - `#pragma omp for`
  - Sections/section constructs
    - `#pragma omp sections`
  - Single construct
    - `#pragma omp single`
- Two directives to be studied
  - **Do/for**: concurrent loop iterations
    - Shares iterations of a loop across the group (data parallelism")
    - Partitions parallel iterations across threads
    - End of for loop: implicit barrier

```
#pragma omp for [clause list]
/* for loop */
```

# Work-Sharing Construct

23

- The parallel and work-sharing (except single) constructs can be combined
  - Loop construct `#pragma omp parallel for`
  - Sections construct `#pragma omp parallel sections`

- Distribute iterations in a parallel region
  - shared clause: All threads can read from and write to a shared variable.
  - private clause: Each thread has a local copy of a private variable

```
#pragma omp parallel for shared(n,a) private(i)  
for (i=0; i <n; i++)  
    a[i] = i + n;
```

- The **mapping between iterations and threads** can be controlled by the **schedule clause**.



# SPMD vs. worksharing

25

- A parallel construct by itself creates a Single Program Multiple Data.
  - each thread redundantly executes the same code.
- Work-sharing in OpenMP

```
for(i=0; i<N; i++) { a[i] = a[i] + b[i];}
```

- OpenMP shortcut: Put the “parallel” and the worksharing directive on the same line

```
#pragma omp parallel for
```

OpenMP parallel  
region work-sharing  
for construct

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    if (id == Nthrds-1) iend = N;
    for(i=istart; i<iend; i++) { a[i] = a[i] + b[i];}
}
```

```
#pragma omp parallel
#pragma omp for
for(i=0; i<N; i++)
{ a[i] = a[i] + b[i];}
```

# Example Using for

26

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
int main()
{
    int nthreads, tid;

    omp_set_num_threads(3);
    #pragma omp parallel private(tid)
    {
        int i;
        tid = omp_get_thread_num();
```

```
        printf("Hello world from (%d)\n", tid);
        #pragma omp for
        for(i = 0; i <=4; i++)
        {
            printf("Iteration %d by %d\n", i, tid);
        }
    }
}
```

Write openMP code to add two vectors using parallel for

## 1. • Two work-sharing loops in one parallel region

```
#pragma omp parallel shared(n,a,b) private(i)
{
  #pragma omp for
  for (i=0; i<n; i++) a[i] = i+1;
  // there is an implied barrier
  #pragma omp for
  for (i=0; i<n; i++) b[i] = 2 * a[i];
}
```

# schedule clause

28

- Defines schedules for OpenMP loops.
- a specification of how iterations of associated loops are divided into contiguous non-empty subsets
  - We call each of the contiguous non-empty subsets a chunk and how these chunks are distributed to threads of the team
- The size of a chunk, denoted as `chunk_size` must be a positive integer

```
#pragma omp parallel for schedule([modifier [modifier]:]kind[,chunk_size])
```

provides a hint for how iterations of the corresponding OpenMP loop should be assigned to threads

## Schedule Types

- static
- dynamic
- guided
- auto
- runtime

# schedule clause

29

- Static
  - Loop iterations are divided into pieces of size chunk
  - If chunk is not specified, the iteration are evenly ) divided contiguously among the threads

Static scheduling • 16 iterations, 4 threads

Thread	0	1	2	3
<i>no chunk*</i>	1-4	5-8	9-12	13-16
<i>chunk = 2</i>	1-2 9-10	3-4 11-12	5-6 13-14	7-8 15-16

# Schedule : Static example

```
int main()
{
    #pragma omp parallel for schedule(static, 3)    for (int i = 0; i < 20; i++)
    {
        printf("Thread %d is running number %d\n", omp_get_thread_num(), i);
    }
    return 0;
}
```

```
int main()
{
    omp_set_num_threads(4);
    #pragma omp parallel for schedule(static, 3)    for (int i = 0; i < 20; i++)
    {
        printf("Thread %d is running number %d\n", omp_get_thread_num(), i);
    }
    return 0;
}
```

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19

Thread 0

Thread 1

Thread 2

Thread 3

Thread 0

Thread 1

Thread 2

## schedule clause

- OpenMP will still split task into `iter_size/chunk_size` chunks, but distribute trunks to threads dynamically without any specific order

`#pragma omp parallel for schedule(dynamic, 1)` is equivalent to `#pragma omp parallel for schedule(dynamic)`

`#pragma omp parallel for schedule(dynamic, chunk-size)`

The dynamic scheduling type is appropriate when the iterations require different computational costs

# schedule clause

32

- **Guided** : Chunk size is dynamic while using guided method, the size of a chunk is proportional to the number of unassigned iterations divided by the number of the threads. Size of each successive chunks is decreasing
  - $\text{chunk size} = \max((\text{num\_of\_iterations remaining} / 2 * \text{num\_of\_threads}), n)$
- **runtime** The scheduling decision is deferred until runtime by the environment variable OMP\_SCHEDULE. It is illegal to specify a chunk size for this clause
- **auto** The scheduling decision is made by the compiler and/or runtime system



- Compiler forms a single loop and then parallelizes this

```
#pragma omp parallel for collapse (2)
```

```
for(i=0;i< N; i++)
```

```
{
```

```
    for(j=0;j< M; j++)
```

```
    {
```

```
        foo(A,i,j);
```

```
    }
```

```
}
```

# Work sharing : Section Directive

- One thread executes one section.  
Each section is executed exactly once.

- **#pragma omp single** (Designated section is executed by single thread only. )
- **#pragma omp master** : Similar to single, but code block will be executed by the master thread only.
- **#pragma omp critical**:

```
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
        x_calculation();
    #pragma omp section
        y_calculation();
    #pragma omp section
        z_calculation();
}
```

- `#pragma omp parallel sections {`
- `#pragma omp section`
- `funcA();`
- `#pragma omp section`
- `funcB(); } /*-- End of parallel region --*/`

# Data Clauses

36

- **Shared clause :**
  - **Syntax:** `shared (item-list)`.
  - Each thread can freely read and modify its value
- **Private Clause**
  - Specifies data that will be replicated so that each thread has a local copy
    - **private clause** (*item-list*)

`firstprivate` and `lastprivate` are just special cases of `private`. The first one results in bringing in values from the outside context into the parallel region while the second one transfers values from the parallel region to the outside context

# Parallel construct

37

- Shared-memory programming model
- Variables are shared by default.
  - **Shared:**
    - All variables visible upon entry of the construct
    - Static variables
  - **Private:**
    - Variables declared within a parallel region
    - (Stack) variables in functions called from within a parallel region
- General rules for data-sharing clauses
  - Clauses default, **private**, **shared**, **firstprivate** allow changing the default behavior
  - Syntax : keyword and a comma-separated list. For instance: `private(a,b)`.

```
int N = 10;
int main( void )
{
    double array[N];
    #pragma omp parallel
    {
        int i, myid;
        double thread_array[N];
        for ( i = 0; i < N; i++ )
            thread_array[i] = myid * array[i];
        function( thread_array );
    }
}

double function( double arg )
{
    static int cnt;
    double local_array[N];
    [...]
}
```

# Firstprivate clause

```
#include <stdio.h>
#include <omp.h>

int main (void)
{
    int i = 10;

    #pragma omp parallel private(i)
    {
        printf("thread %d: i = %d\n", omp_get_thread_num(), i);
        i = 1000 + omp_get_thread_num();
    }

    printf("i = %d\n", i);

    return 0;
}
```

When i is made firstprivate

```
thread 2: i = 10
thread 0: i = 10
thread 3: i = 10
thread 1: i = 10
i = 10
```

```
thread 0: i = 0
thread 3: i = 32717
thread 1: i = 32717
thread 2: i = 1
i = 10
```

(another run of the same program)

```
thread 2: i = 1
thread 1: i = 1
thread 0: i = 0
thread 3: i = 32657
i = 10
```

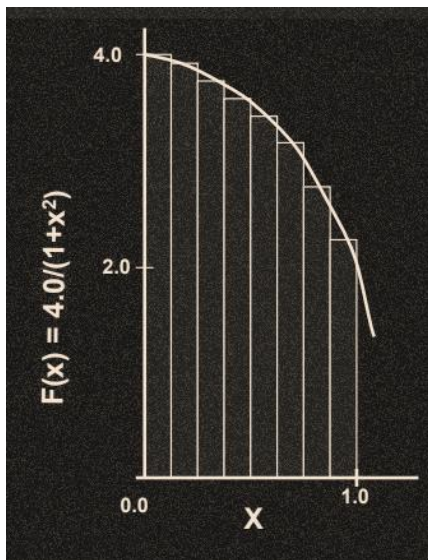
# Lastprivate clause

39

```
#pragma omp parallel for private(i) lastprivate(a)
for (i=0; i<n; i++) {
    a = i+1;
    printf("Thread %d has a value of a = %d for i = %d\n", omp_get_thread_num(),a,i);
} /*-- End of parallel for --*/
printf("After parallel for: i = %d , a = %d\n", i, a);
```

Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$



$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

```
static long num_steps = 100000;
```

```
double step;
```

```
void main ()
```

```
{
```

```
    int i; double x, pi, sum = 0.0;
```

```
    step = 1.0/(double) num_steps;
```

```
    for (i=0; i< num_steps; i++)
```

```
    {
```

```
        x= (i+0.5)*step; sum = sum + 4.0/(1.0+x*x);
```

```
    }
```

```
    pi = step * sum;
```

```
}
```

**Serial PI Program**

Numerical Integration Example



# Homework

41

1. Create a parallel version of the pi program using a parallel construct. ???

```
#include <stdio.h>
#include <stdlib.h>
#define NUM_STEPS 10000
int main( void )
{
    int i;
    double sum = 0.0, pi, x_i;
    double step = 1.0/NUM_STEPS;
    for ( i = 0; i < NUM_STEPS; i++ ) {
        x_i = (i + 0.5) * step;
        sum = sum + 4.0 / (1.0 + x_i * x_i);
    }
    pi = sum * step;
    printf("Pi: %.15e\n", pi);
    return 0;
}
```

Hints:

```
#pragma omp parallel
num_threads(...)
omp_set_num_threads(...)
omp_get_num_threads(...)
omp_get_thread_num(...)
```

Challenges:

- split iterations of the loop among threads
- create an accumulator for each thread to hold partial sums, which can later be combined to generate the global sum

```
int main()
{
    int b[3];
    char *cptr;
    int i;

    cptr = malloc(1);
    #pragma omp parallel for
    for(i=0; i<3; i++)
        b[i]=i;
}
```

1. List out the private and shared variable in following code

1. We want each thread's private copy of array element  $x[0]$  to inherit the value that the shared variable was assigned in the master thread

```
x[0] = 1.0;
for(i=0; i < n; i++){
    for(j=1; j<4; j++)
        x[j]=g(i, x[j-1]);
    answer[i]=x[1]-x[3];
}
```

`#pragma omp parallel for private (j) firstprivate (x)`

- How do we decide which variables should be shared and which private?
  - Loop indices - private –
  - Loop temporaries - private –
  - Read-only variables - shared –
  - Main arrays - shared



# Example

45

- Write OpenMP program to calculate the average of an array of integers.
  - The program should efficiently distribute the workload among multiple threads, and you must handle critical sections appropriately to avoid data races and ensure accurate results.

# Solution:

46

```
// Initialize the array with values
for (int i = 0; i < ARRAY_SIZE; ++i) {
    array[i] = i + 1;
}
```

```
#pragma omp parallel num_threads(num_threads) shared(sum)
{
    double local_sum = 0.0;
```

```
#pragma omp for
for (int i = 0; i < ARRAY_SIZE; ++i) {
    local_sum += array[i];
}
```

```
#pragma omp critical
{
    sum += local_sum;
}
}
```

```
double average = sum /
ARRAY_SIZE;
```

```
printf("Average: %f\n", average);
```

- Design a parallel program in C using OpenMP to perform a linear search on an array.
- The goal is to
  - efficiently search for a specific target
  - Handle critical sections appropriately to ensure the correct identification

```
#pragma omp parallel for
for (int i = 0; i < size; ++i) {
    if (arr[i] == target) {
        #pragma omp critical
        {
            result = i; // Set result only once using a critical section
        }
    }
}
```



# Reduction

49

- Reduction (operator: variable list):
  - Combining of local copies of a variable in different threads into a single copy at the master when threads exit.
  - Variables in variable list are implicitly private to threads.
- Operators: +, \*, -, &, |, ^, &&, and || – Usage

```
#pragma omp parallel reduction(+: sums) num_threads(4)
{ /* compute local sums in each thread }
```

# Example

50

```
sum = 0;
for (int i = 0; i < 10; i++)
{
    sum += a[i]
}
```

#pragma omp parallel for shared(sum, a) **reduction(+: sum)**

```
sum = 0;
#pragma omp parallel for shared(sum, a) reduction(+: sum)
for ( int i = 0; i < 10; i++) {
    sum += a[i] }
```

```
for( int i = 0; i < numSubdivisions; i++ )  
{ double x = A + dx * (float) i;  
  double f = Function( x );  
  sum += f;  
}
```

There are Three Ways to Make the Summing Work Correctly:

```
#pragma omp atomic  
sum += f;
```

```
#pragma omp critical  
sum += f;
```

```
#pragma omp parallel for shared(dx),reduction(+:sum)
```

# Advantages and Disadvantages openMP

52

## Advantages

- Shared address space provides user friendly programming
- Ease of programming
- Data sharing between threads is fast and uniform (low latency)
- Incremental parallelization of sequential code
- Leaves thread management to compiler
- Directly supported by compiler

## Disadvantages

- Internal details are hidden
- Programmer is responsible for specifying synchronization, e.g. locks
- Cannot run across distributed memory
- Performance limited by memory architecture
- Lack of scalability between memory and CPUs
- Requires compiler which supports OpenMP

# Monte Carlo to estimate PI

53

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
int main(int argc, char *argv[])
{
    long int i, count; // count points
    inside unit circle
    long int samples; // number of samples
    double pi;
    unsigned short xi[3] = {1, 5, 177}; //
    random number seed
    double x, y;
    samples = atoi(argv[1]);
    count = 0;
    for(i = 0; i < samples; i++)
    {
```

```
        x = erand48(xi);
        y = erand48(xi);
        if(x*x + y*y <= 1.0) count++;
    }
    pi = 4.0*count/samples;
    printf("Estimate of pi: %7.5f\n", pi);
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include "omp.h"
main(int argc, char *argv[])
{
    int i, count; /* points inside the unit quarter
circle */
    unsigned short xi[3]; /* random number seed */
    int samples; /*samples Number of points to
Generate*/
    double x,y; /* Coordinates of points */
    double pi; /* Estimate of pi */
    samples = atoi(argv[1]);
    #pragma omp parallel
    {
        xi[0] = 1; /* set up the random seed */
        xi[1] = 1;
        xi[2] = omp_get_thread_num();
        count = 0;
```

```
        printf("I am thread %d\n", xi[2]);
        #pragma omp for firstprivate(xi) private(x,y)
        reduction(+:count)
        for (i = 0; i < samples; i++) {
            x = erand48(xi);
            y = erand48(xi);
            if (x*x + y*y <= 1.0) count++; }
    }
    pi = 4.0 * (double)count / (double)samples;

    printf("Count = %d, Samples = %d, Estimate of pi:
%.5f\n", count, samples, pi);
}
```


# Mid-QP

55

```
for (i=0; i<m; i++)  
{  
    a[i] = 0.0;  
    for (j=0; j<n; j++)  
        a[i] += b[i][j]*c[j];  
}
```

```
#pragma omp parallel for default(none) \  
    private(i,j) shared(m,n,a,b,c)
```

```
for (i=0; i<m; i++)  
{  
    a[i] = 0.0;  
    for (j=0; j<n; j++)  
        a[i] += b[i][j]*c[j];  
}
```



# Tutorial 0

56

1. You are given an array `shared_array` of size `ARRAY_SIZE`. Multiple threads need to concurrently increment each element of the array.
2. Design a parallel program using OpenMP to achieve this, ensuring that the increments are performed atomically to avoid race conditions.



# Tutorial Exercise 1

57

- The SAXPY program is to add a scalar multiple of a real vector to another real vector:
- $s = a * x + y$ .
- Provided a serial SAXPY code, parallelize it using OpenMP directives.
- Compare the performance between serial and OpenMP codes.

```
for { i = 0; i < n; i++ }  
{  
    y[i] = a * x[i] + y[i];  
}
```

- Check total wall clock execution time versus thread numbers:
  - `export OMP_NUM_THREADS=1`
  - `time ./saxpy`
  - `export OMP_NUM_THREADS=2`
  - `time ./saxpy`
  - `export OMP_NUM_THREADS=4`
  - `time ./saxpy`
  - `export OMP_NUM_THREADS=8`
  - `time ./saxpy`

# Tutorial Exercise 2

59

## 1. Estimating the value of Pi using Monte Carlo


1. computational algorithms that rely on repeated random sampling to obtain numerical results

### Estimation of Pi

We know that area of the square is  $4r^2$  unit sq while that of circle is  $\pi r^2$ .

The ratio of these two areas is as follows :  $\pi/4$

1. Generate a random point  $(x, y)$  inside a square of side 2 centered at the origin.
2. Determine whether the point falls inside the unit circle inscribed in the square by checking whether  $x^2 + y^2 \leq 1$ .
3. Repeat steps 1 and 2 for a large number of points (e.g.,  $10^7$ ).
4. Calculate the ratio of the number of points that fell inside the circle to the total number of points generated.
5. Multiply the ratio by 4 to estimate the value of  $\pi$ .

MC Approximation of  $\pi = 3.14616$ 

# Monte Carlo to estimate PI

61

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
int main(int argc, char *argv[])
{
    long int i, count; // count points
    inside unit circle
    long int samples; // number of samples
    double pi;
    unsigned short xi[3] = {1, 5, 177}; //
    random number seed
    double x, y;
    samples = atoi(argv[1]);
    count = 0;
    for(i = 0; i < samples; i++)
    {
```

```
        x = erand48(xi);
        y = erand48(xi);
        if(x*x + y*y <= 1.0) count++;
    }
    pi = 4.0*count/samples;
    printf("Estimate of pi: %7.5f\n", pi);
}
```

- Provided a serial code, parallelize it using OpenMP directives.
- Compare the performance between serial and OpenMP codes

# References

63

1. <http://openmp.org> 2.
2. <https://computing.llnl.gov/tutorials/openMP> 3.
3. [http://www.openmp.org/mp-documents/OpenMP4.0RC1\\_final.pdf](http://www.openmp.org/mp-documents/OpenMP4.0RC1_final.pdf)
4. Michael J. Quinn, Parallel Programming in C with MPI and OpenMP, Mc Graw Hill, 2003.